

Deep Learning-based Leak Detection for Smart Water Metering with Consumption Forecasting

Problem Identified

DENR airs concern over water leaks in houses, buildings

Bella Cariaso - The Philippine Star ①

May 9, 2024 | 12:00am



Almost 1.7 cubic meters per second of water lost due to pipe leaks —Manila Water

By GISELLE OMBAY, GMA Integrated News

Published April 3, 2023 10:15am

Updated April 3, 2023 11:36am



Problem Statement

Water pipeline leaks cause significant losses in urban systems, yet traditional detection methods are often manual, inaccurate, and inefficient. This project proposes a deep learning-based approach that combines automated leak detection and water consumption forecasting using images and numerical data. Our goal is to enhance detection accuracy, and support smarter, real-time water management.





Dataset Chosed

OBJECT DETECTION

Pipeline Leak Prediction (Roboflow)

<https://universe.roboflow.com/gas-leak/pipeline-leak-prediction>

Contents:

images of pipelines that has leaks and no leaks

FOR LSTM


A synthetic dataset made of water consumption consisting 10 thousands of sets of it.

Methodology

Dual-Model Approach: Object Detection
and Time Series Analysis




Leak Detection Model



Detectron2

RetinaNet

Consumption Forecasting Model



**TIME SERIES
PREDICTION
WITH
GRU**

TSN Time Sensitive™
Networking Powered by
IEEE 802

Dataset prep for YOLOv11 Model

Pipeline Leak Prediction Dataset (from Roboflow)					
SET	Total Images	Leak	No leak	Water	Crack
Training	2739	1099	623	847	170
Validation	488	193	91	168	36
Total	3227	1292	714	1015	206



Data Augmentation
already applied

85% training
15% validation

Preprocessed via
Roboflow

Training of Objection Detection Models



- enhance feature extraction + object localization
- high inference speed and improved detection precision



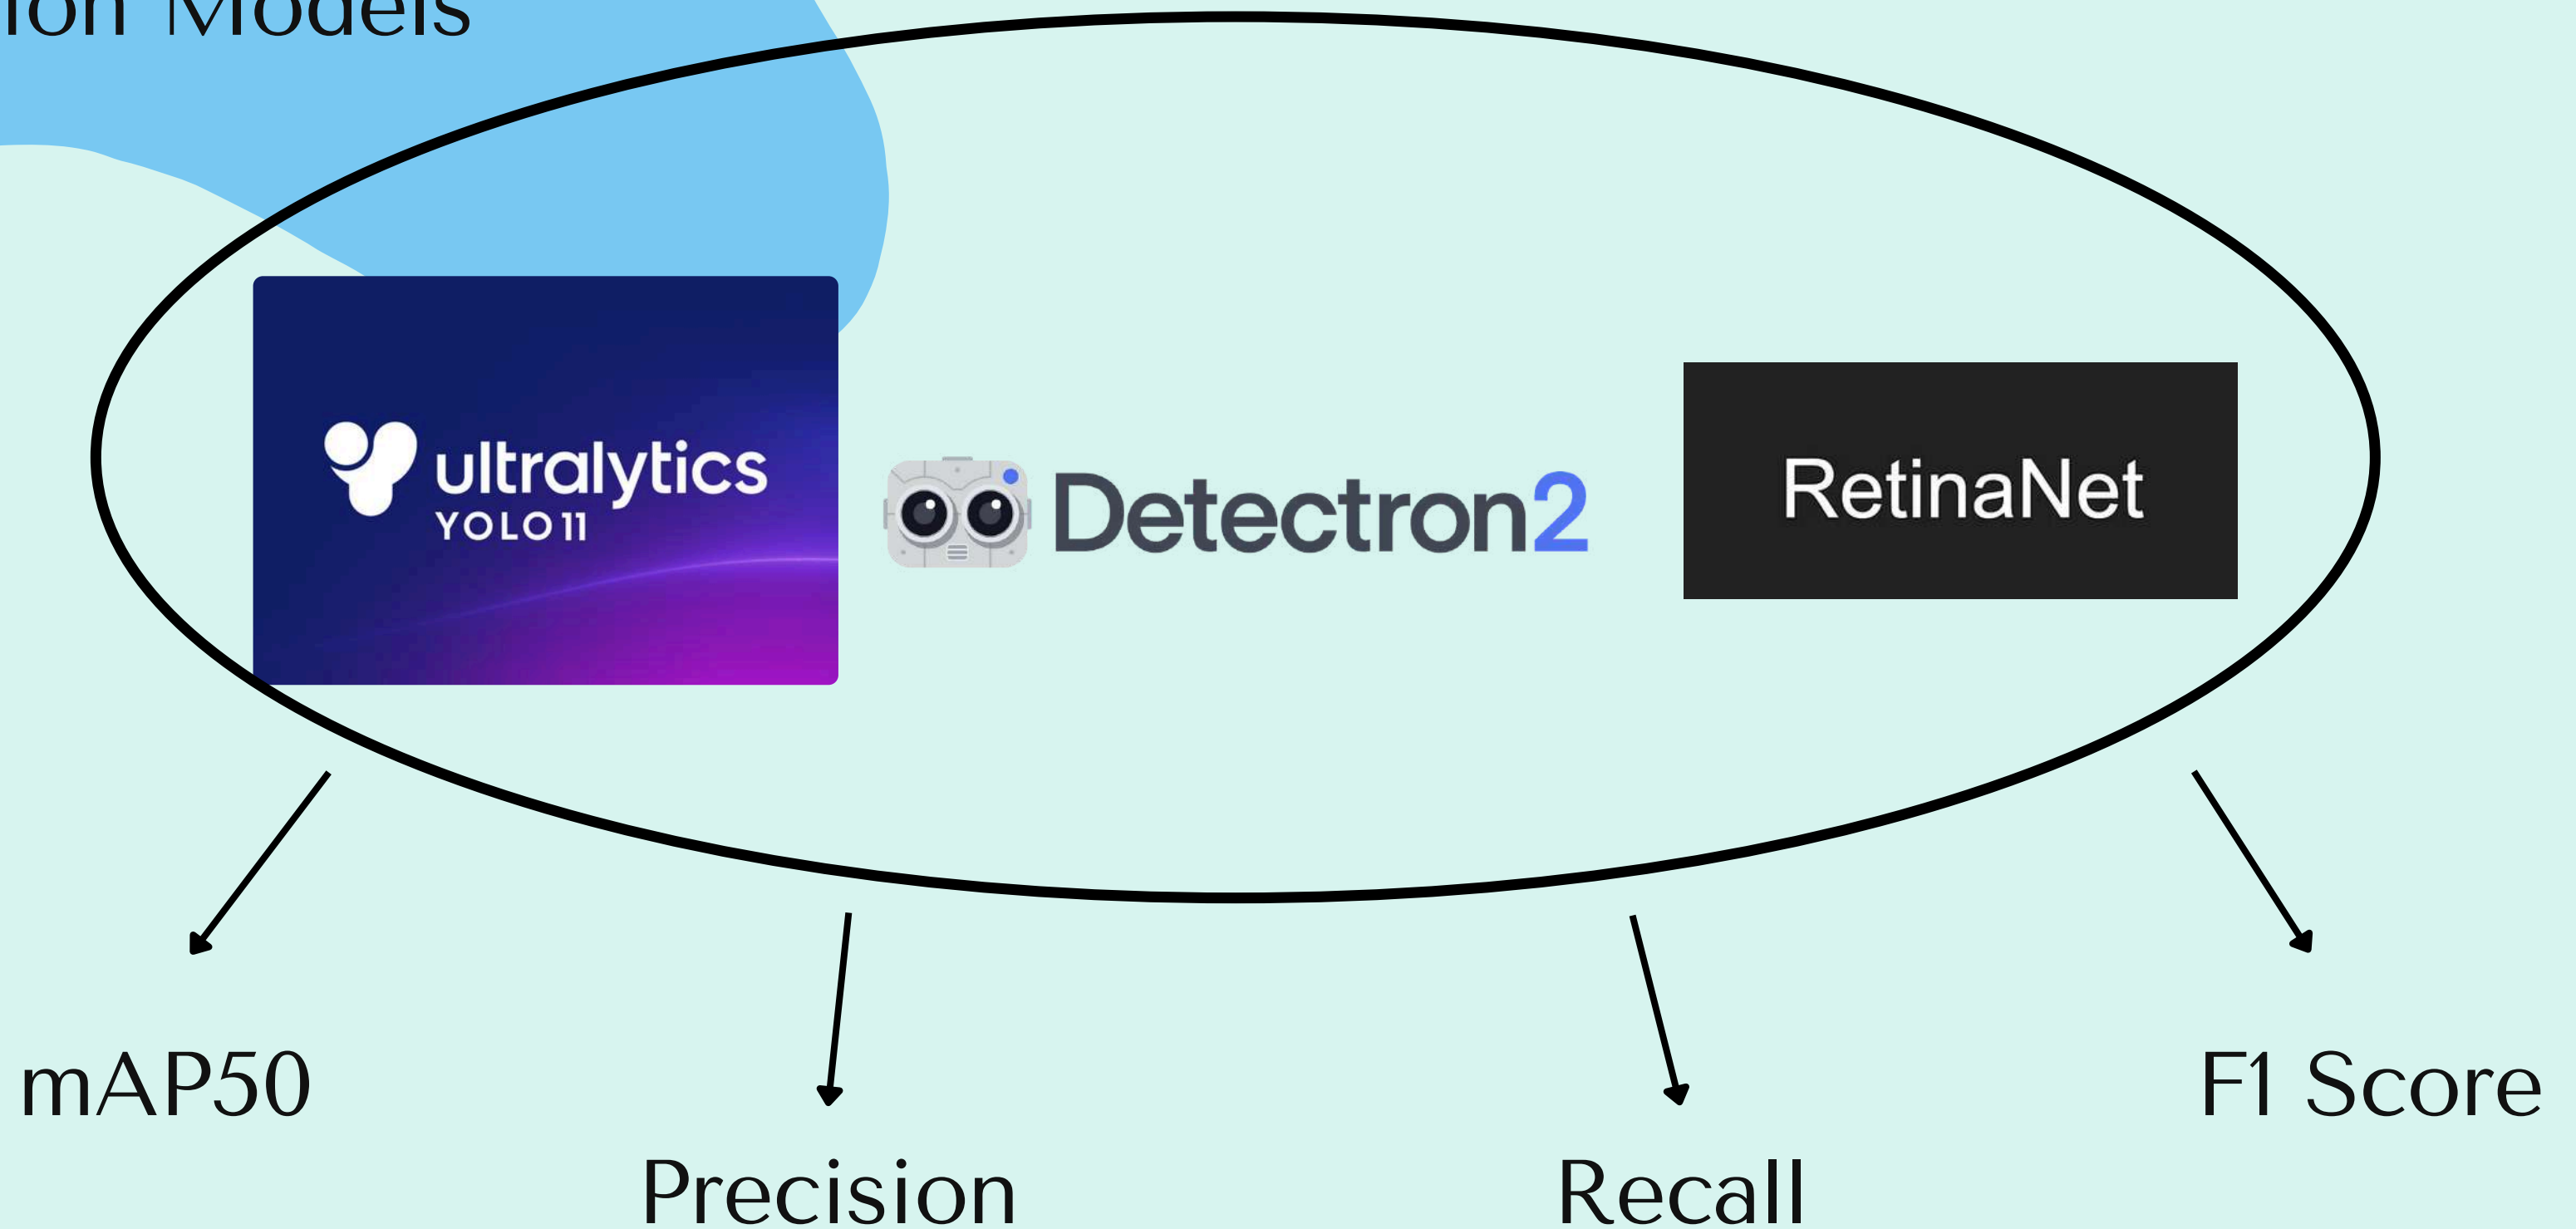
Detectron2

- fast training on single or multiple GPU servers and includes high-quality implementations of numerous advanced models like Mask R-CNN, Cascade R-CNN, and Panoptic FPN.

RetinaNet

- Addresses Class Imbalance in One-Stage Detectors
- Balances Speed and Accuracy

Evaluating the performance of Leak Detection Models



Deployment of Leak Detection Model in Streamlit



Detectron2





 Streamlit

RetinaNet


Results and Discussion




Leak Detection Model
Performance Evaluation



Performance Metrics of YOLOv11 model



- peaking at 0.81 around a confidence threshold of ~ 0.527 . This shows good indication that the yolo model performs reliably when it is reasonably confident





- Precision and recall trade-offs are well balanced for most classes except “no leak.”
- An mAP@0.5 of .860 indicates good object localization and classification confidence, particularly for leak-related targets.

Performance Metrics
of YOLOv11 model


YOLOv11 Model Performance Evaluation				
Class	Images	Instances	R	mAP50
all	502	492	0.767	0.859
crack	36	36	0.832	0.946
leak	193	199	0.899	0.942
no leak	91	92	0.554	0.663
water	168	169	0.783	0.885

1st test of Detectron2 model

```
cfg = get_cfg()
cfg.merge_from_file(model_zoo.get_config_file("COCO-Detection/faster_rcnn_R_50_FPN_3x.yaml"))
cfg.DATASETS.TRAIN = ("leak_dataset_train",)
cfg.DATASETS.TEST = ("leak_dataset_val",)
cfg.DATALOADER.NUM_WORKERS = 2
cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url("COCO-Detection/faster_rcnn_R_50_FPN_3x.yaml")
cfg.SOLVER.IMS_PER_BATCH = 2
cfg.SOLVER.BASE_LR = 0.00125
cfg.SOLVER.MAX_ITER = 2000
cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 128
```



1st test of Detectron2 model



1st test of Detectron2 model



```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.344
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.497
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.370
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.212
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.348
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.534
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.622
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.631
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.268
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.654
[05/22 16:31:55 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | AP1 |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 34.432 | 49.703 | 36.961 | nan | 21.236 | 34.785 |
[05/22 16:31:55 d2.evaluation.coco_evaluation]: Some metrics cannot be computed and is shown as NaN.
[05/22 16:31:55 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| leakage | nan | crack | 42.730 | leak | 48.876 |
| no leak | 10.821 | water | 35.301 | | |
OrderedDict([('bbox',
               {'AP': 34.43207423228552,
                'AP50': 49.70307085589206,
                'AP75': 36.961087166617936,
                'APs': nan,
                'APm': 21.23628686307657,
                'AP1': 34.78472621458232,
                'AP-leakage': nan,
                'AP-crack': 42.73049901827959,
                'AP-leak': 48.876384873901344,
                'AP-no leak': 10.820891749531853,
                'AP-water': 35.3005212874293})])
```


2nd test of Detectron2 model


```
cfg = get_cfg()
cfg.merge_from_file(model_zoo.get_config_file("COCO-Detection/faster_rcnn_R_50_FPN_3x.yaml"))

cfg.DATASETS.TRAIN = ("leak_dataset_train",)
cfg.DATASETS.TEST = ("leak_dataset_val",)
cfg.DATALOADER.NUM_WORKERS = 2

cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url("COCO-Detection/faster_rcnn_R_50_FPN_3x.yaml")
cfg.SOLVER.IMS_PER_BATCH = 4
cfg.SOLVER.BASE_LR = 0.00125
cfg.SOLVER.MAX_ITER = 5000
cfg.SOLVER.STEPS = []
cfg.TEST.EVAL_PERIOD = 500
cfg.SOLVER.OPTIMIZER = "ADAMW"
cfg.SOLVER.WEIGHT_DECAY = 0.0005
cfg.SOLVER.AMSGRAD = True
cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 32
cfg.MODEL.ROI_HEADS.SCORE_THRESH_TEST = 0.5
```



2nd test of Dectectron 2 model



2nd test of Detelectron2 model

```
[05/22 17:53:23 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.02 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.571
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.742
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.608
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.375
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.592
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.651
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.700
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.700
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.535
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.716
[05/22 17:53:23 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
|  AP   |  AP50  |  AP75  |  APs   |  APm   |  APl   |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 57.051 | 74.164 | 60.793 | nan    | 37.521 | 59.171 |
[05/22 17:53:23 d2.evaluation.coco_evaluation]: Some metrics cannot be computed and is shown as NaN.
[05/22 17:53:23 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP   | category | AP   | category | AP   |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| leakage  | nan   | crack    | 63.591 | leak     | 67.730 |
| no leak  | 38.043 | water    | 58.841 |          |        |
```




1st test on Retinanet Model

```
# Load configuration from a RetinaNet config file
cfg.merge_from_file(model_zoo.get_config_file("COCO-Detection/retinanet_R_50_FPN_3x.yaml"))

cfg.DATASETS.TRAIN = ("leak_dataset_train",)
cfg.DATASETS.TEST = ("leak_dataset_val",)
cfg.DATALOADER.NUM_WORKERS = 2

# Load weights from a RetinaNet checkpoint
cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url("COCO-Detection/retinanet_R_50_FPN_3x.yaml")

cfg.SOLVER.IMS_PER_BATCH = 2
cfg.SOLVER.BASE_LR = 0.00125
cfg.SOLVER.MAX_ITER = 2000
num_classes_from_dataset = len(leak_metadata.thing_classes)
cfg.MODEL.RETINANET.NUM_CLASSES = num_classes_from_dataset
```






1st test on Retinanet Model

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.411
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.580
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.437
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.253
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.416
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.520
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.622
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.633
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.324
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.650
[05/18 18:53:10 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
|  AP   | AP50 | AP75 | APs  | APm  | APl  |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 41.078 | 58.032 | 43.691 | nan  | 25.302 | 41.558 |
[05/18 18:53:10 d2.evaluation.coco_evaluation]: Some metrics cannot be computed and is shown as NaN.
[05/18 18:53:10 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP   | category | AP   | category | AP   |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| leakage  | nan   | crack    | 44.277 | leak     | 57.518 |
| no leak  | 14.504 | water    | 48.013 |          |         |
OrderedDict([('bbox'
```


2nd test of Retinanet model

```
cfg_tuned = get_cfg()
cfg_tuned.merge_from_file(model_zoo.get_config_file("COCO-Detection/faster_rcnn_R_50_FPN_3x.yaml"))
cfg_tuned.DATASETS.TRAIN = ("leak_dataset_train",)
cfg_tuned.DATASETS.TEST = ("leak_dataset_val",)
cfg_tuned.DATALOADER.NUM_WORKERS = 2 # Increased number of workers
cfg_tuned.MODEL.WEIGHTS = model_zoo.get_checkpoint_url("COCO-Detection/faster_rcnn_R_50_FPN_3x.yaml")
cfg_tuned.SOLVER.IMS_PER_BATCH = 4 # Increased batch size
cfg_tuned.SOLVER.BASE_LR = 0.00125 # Increased learning rate
cfg_tuned.SOLVER.MAX_ITER = 5000 # Increased iterations
cfg_tuned.SOLVER.STEPS = (4000, 4500) # Add learning rate decay steps
cfg.SOLVER.OPTIMIZER = "ADAMW"
cfg.TEST.EVAL_PERIOD = 500
cfg.SOLVER.WEIGHT_DECAY = 0.0005
cfg_tuned.SOLVER.GAMMA = 0.1 # Learning rate decay factor
cfg_tuned.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 128 # Increased RoIHead batch size
cfg_tuned.MODEL.ROI_HEADS.NUM_CLASSES = num_classes_from_dataset

os.makedirs(cfg_tuned.OUTPUT_DIR, exist_ok=True)
trainer_tuned = DefaultTrainer(cfg_tuned)
trainer_tuned.resume_or_load(resume=False)
trainer_tuned.train()
```




2nd test of Retinanet model

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.499
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.695
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.547
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.210
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.522
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.608
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.664
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.680
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.347
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.698
[05/18 20:36:21 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
|  AP   |  AP50  |  AP75  |  APs   |  APm   |  APl   |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 49.855 | 69.490 | 54.737 | nan    | 20.981 | 52.226 |
[05/18 20:36:21 d2.evaluation.coco_evaluation]: Some metrics cannot be computed and is shown as NaN.
[05/18 20:36:21 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP   | category | AP   | category | AP   |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| leakage  | nan   | crack    | 51.123 | leak     | 65.168 |
| no leak  | 30.524 | water    | 52.604 |           |         |
```



Conclusion and Recommendation



Which is the best model?


- Highest Detection Accuracy — Achieved an mAP50 of 85.9%, outperforming Detectron2 and RetinaNet.
- Fast and Real-Time — Optimized for low-latency inference, ideal for smart water metering systems.
- Strong Recall on Leaks — Detected nearly 90% of actual leaks, ensuring critical issues aren't missed.
- Robust Generalization — Handled diverse image conditions thanks to built-in data augmentations.
- Stable and Efficient Training — Smooth convergence without overfitting, simplifying deployment.

ultralyti
YOLO11

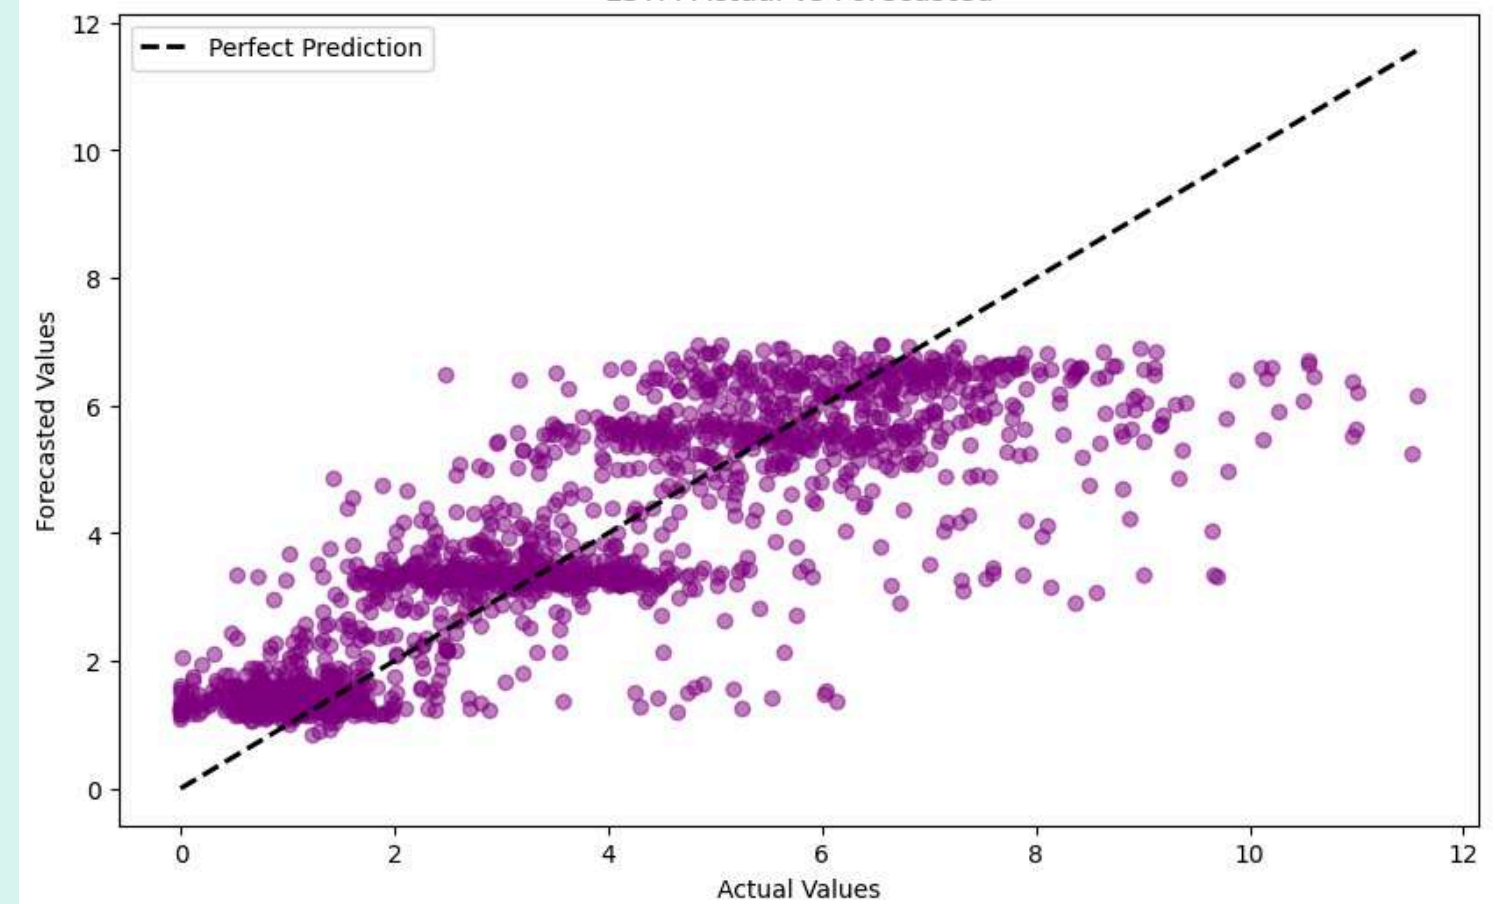
LSTM

```
# Optimized LSTM Model with Dropout and Early Stopping
model_lstm = Sequential([
    Input(shape=(window_size, 1)),
    LSTM(64, return_sequences=True),
    Dropout(0.2),
    LSTM(32),
    Dropout(0.2),
    Dense(1)
])
model_lstm.compile(optimizer='adam', loss='mse')
early_stopping = EarlyStopping(patience=7, restore_best_weights=True)
model_lstm.fit(X_train, y_train, validation_split=0.1, epochs=50, batch_size=32,
               callbacks=[early_stopping], verbose=1)
```

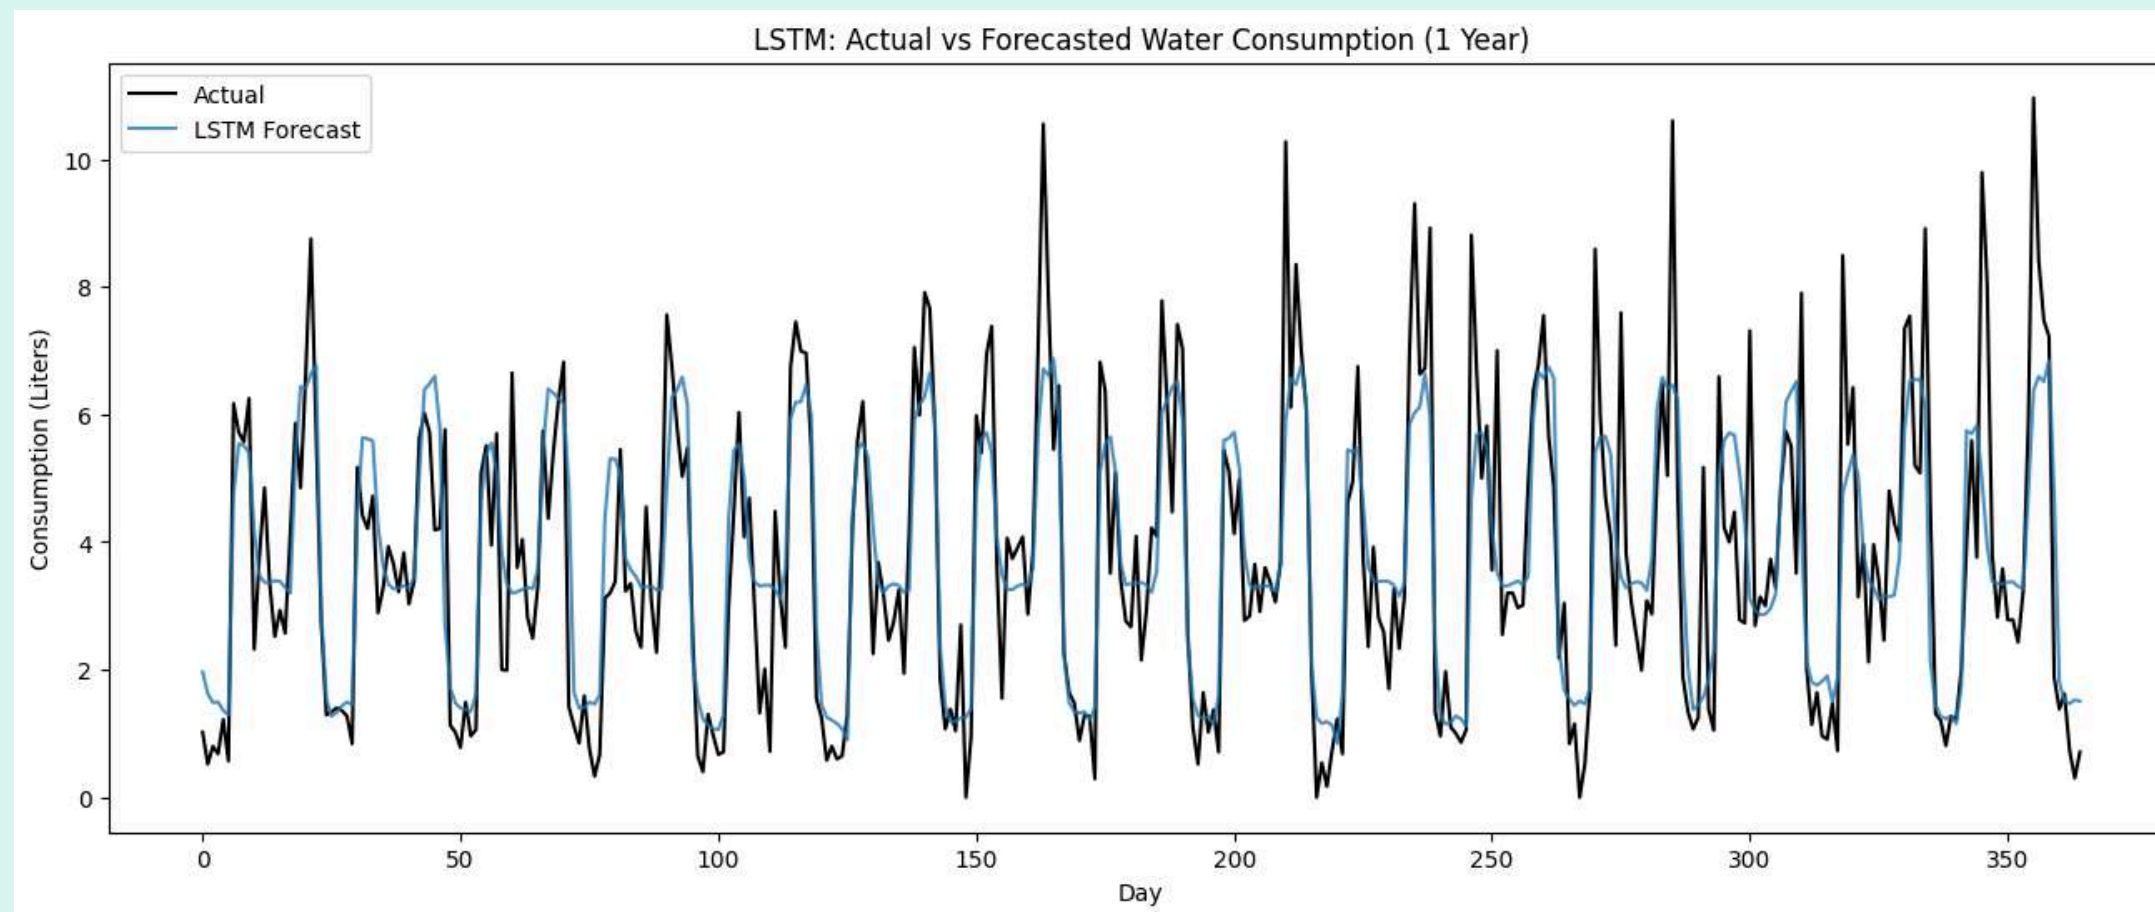
LSTM Forecasting



LSTM Actual vs Forecasted

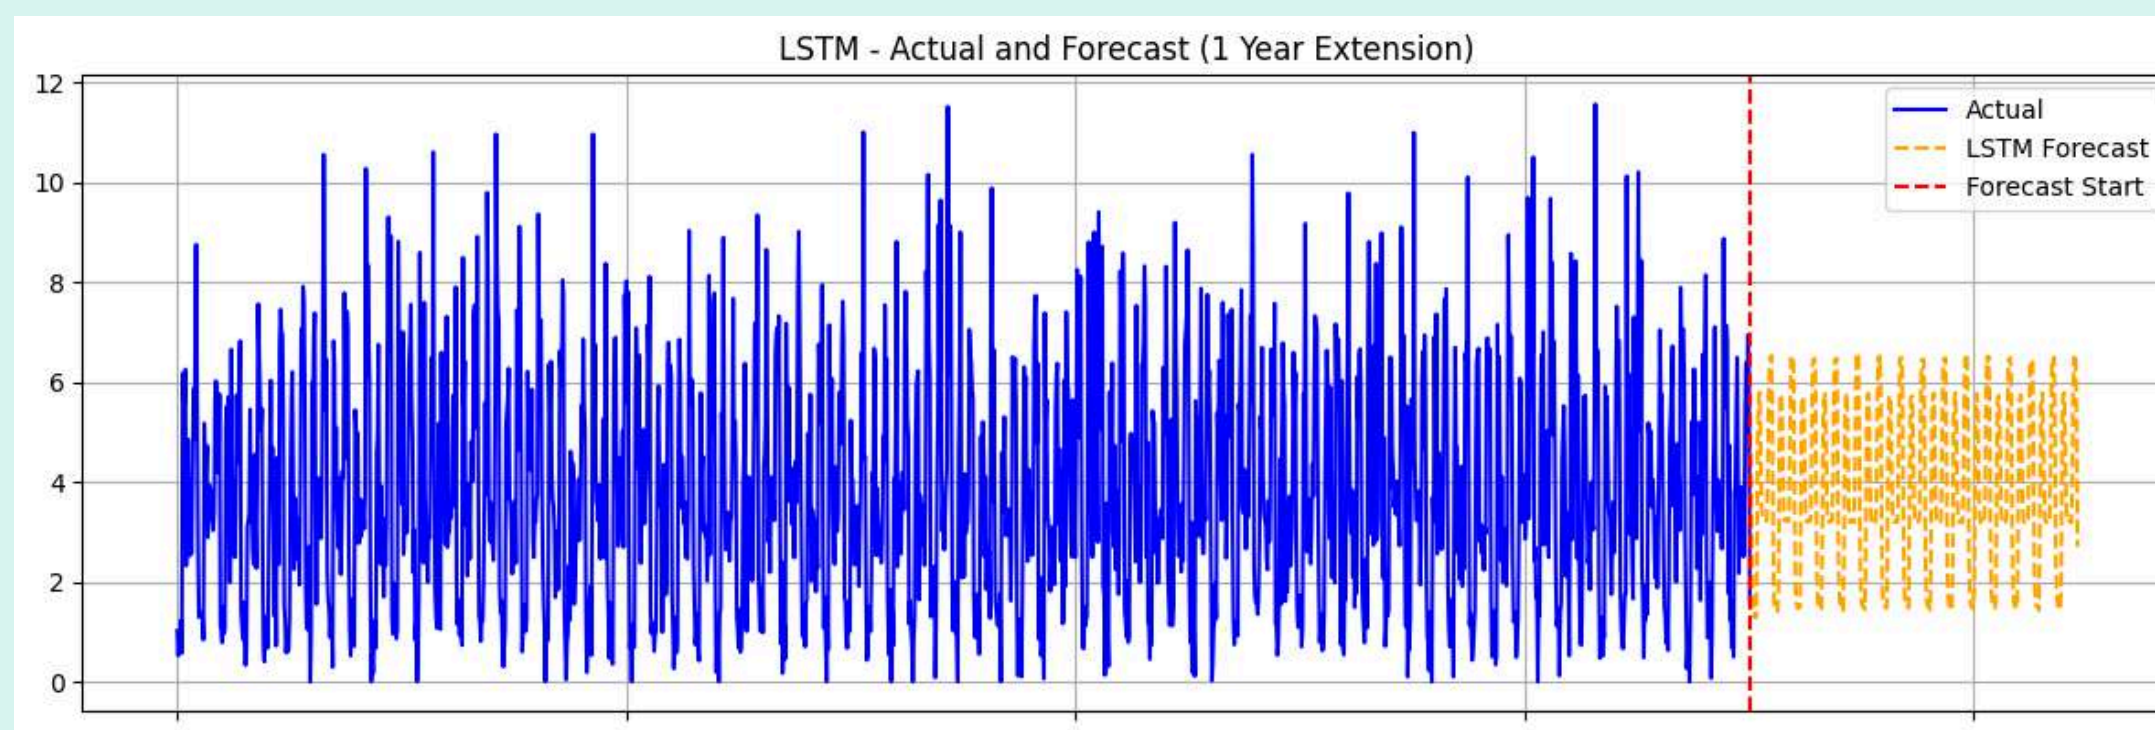


LSTM



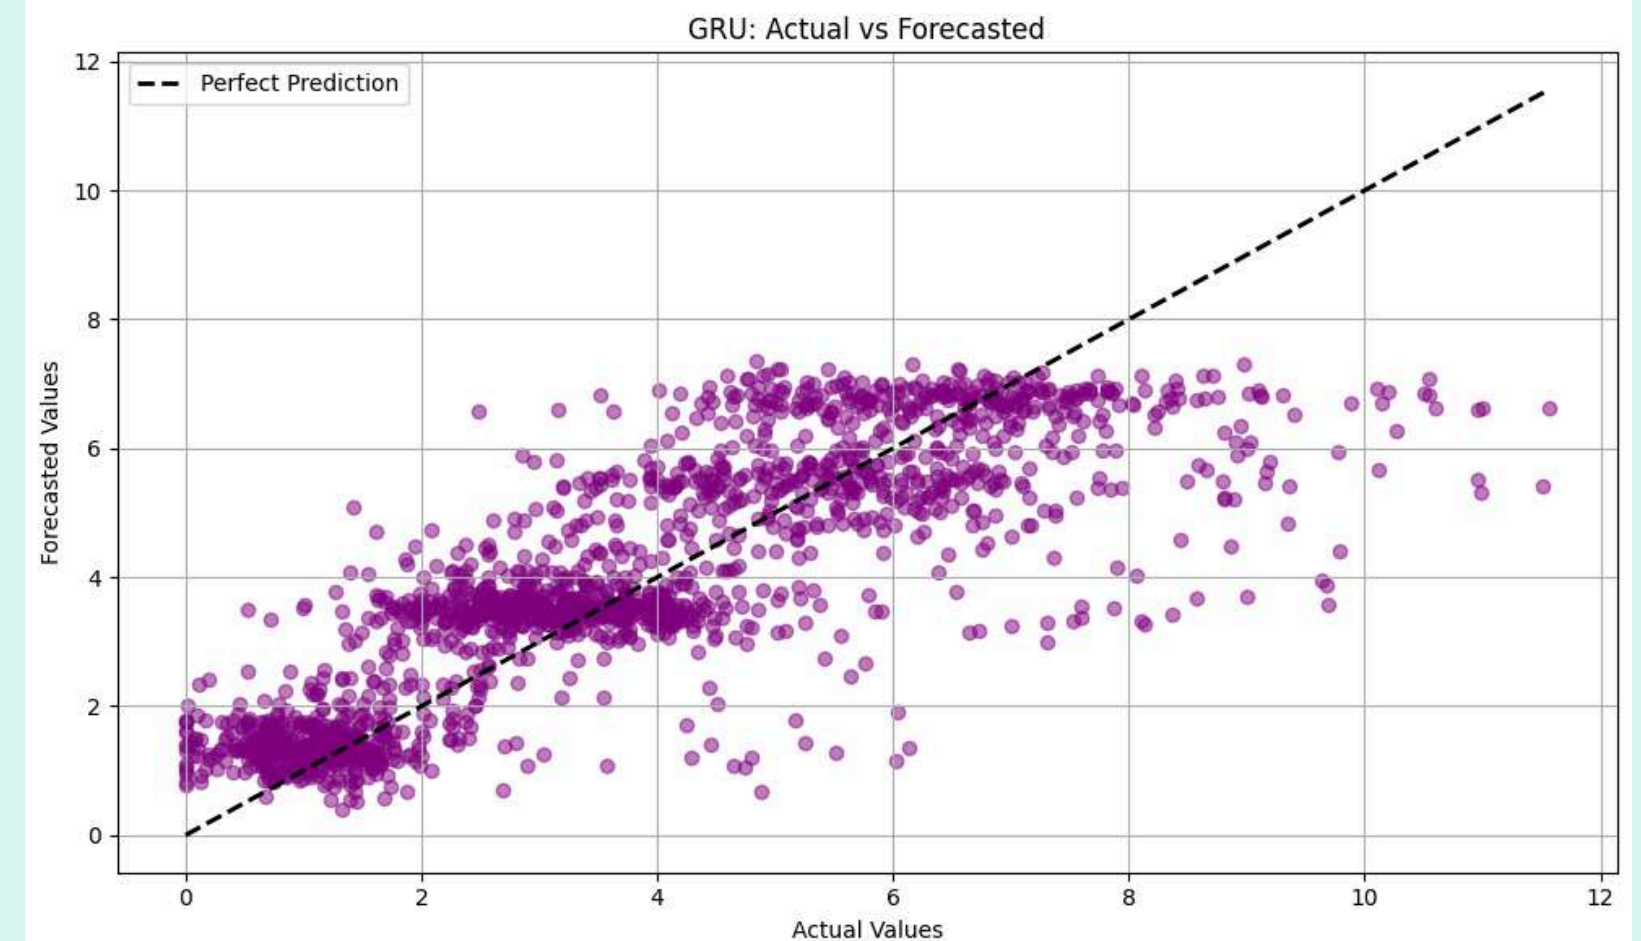
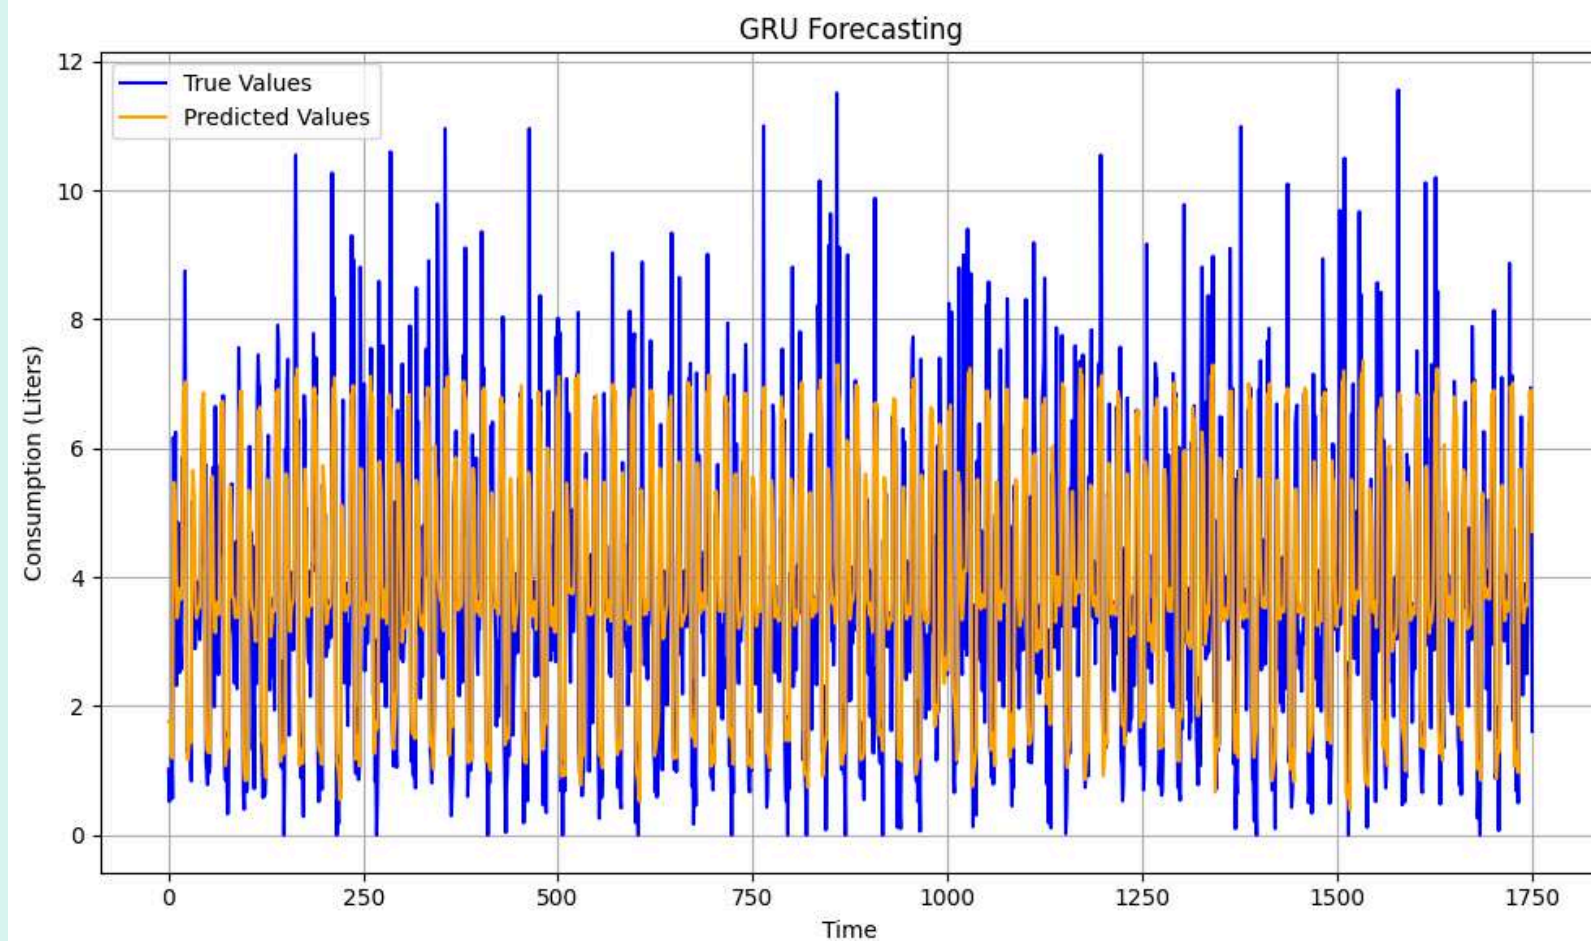
Forecast Behavior: Produces smooth and stable forecasts with clear periodic patterns. Strengths: Good at capturing long-term dependencies and seasonality.

Limitations: May underrepresent short-term fluctuations or sudden changes in consumption.


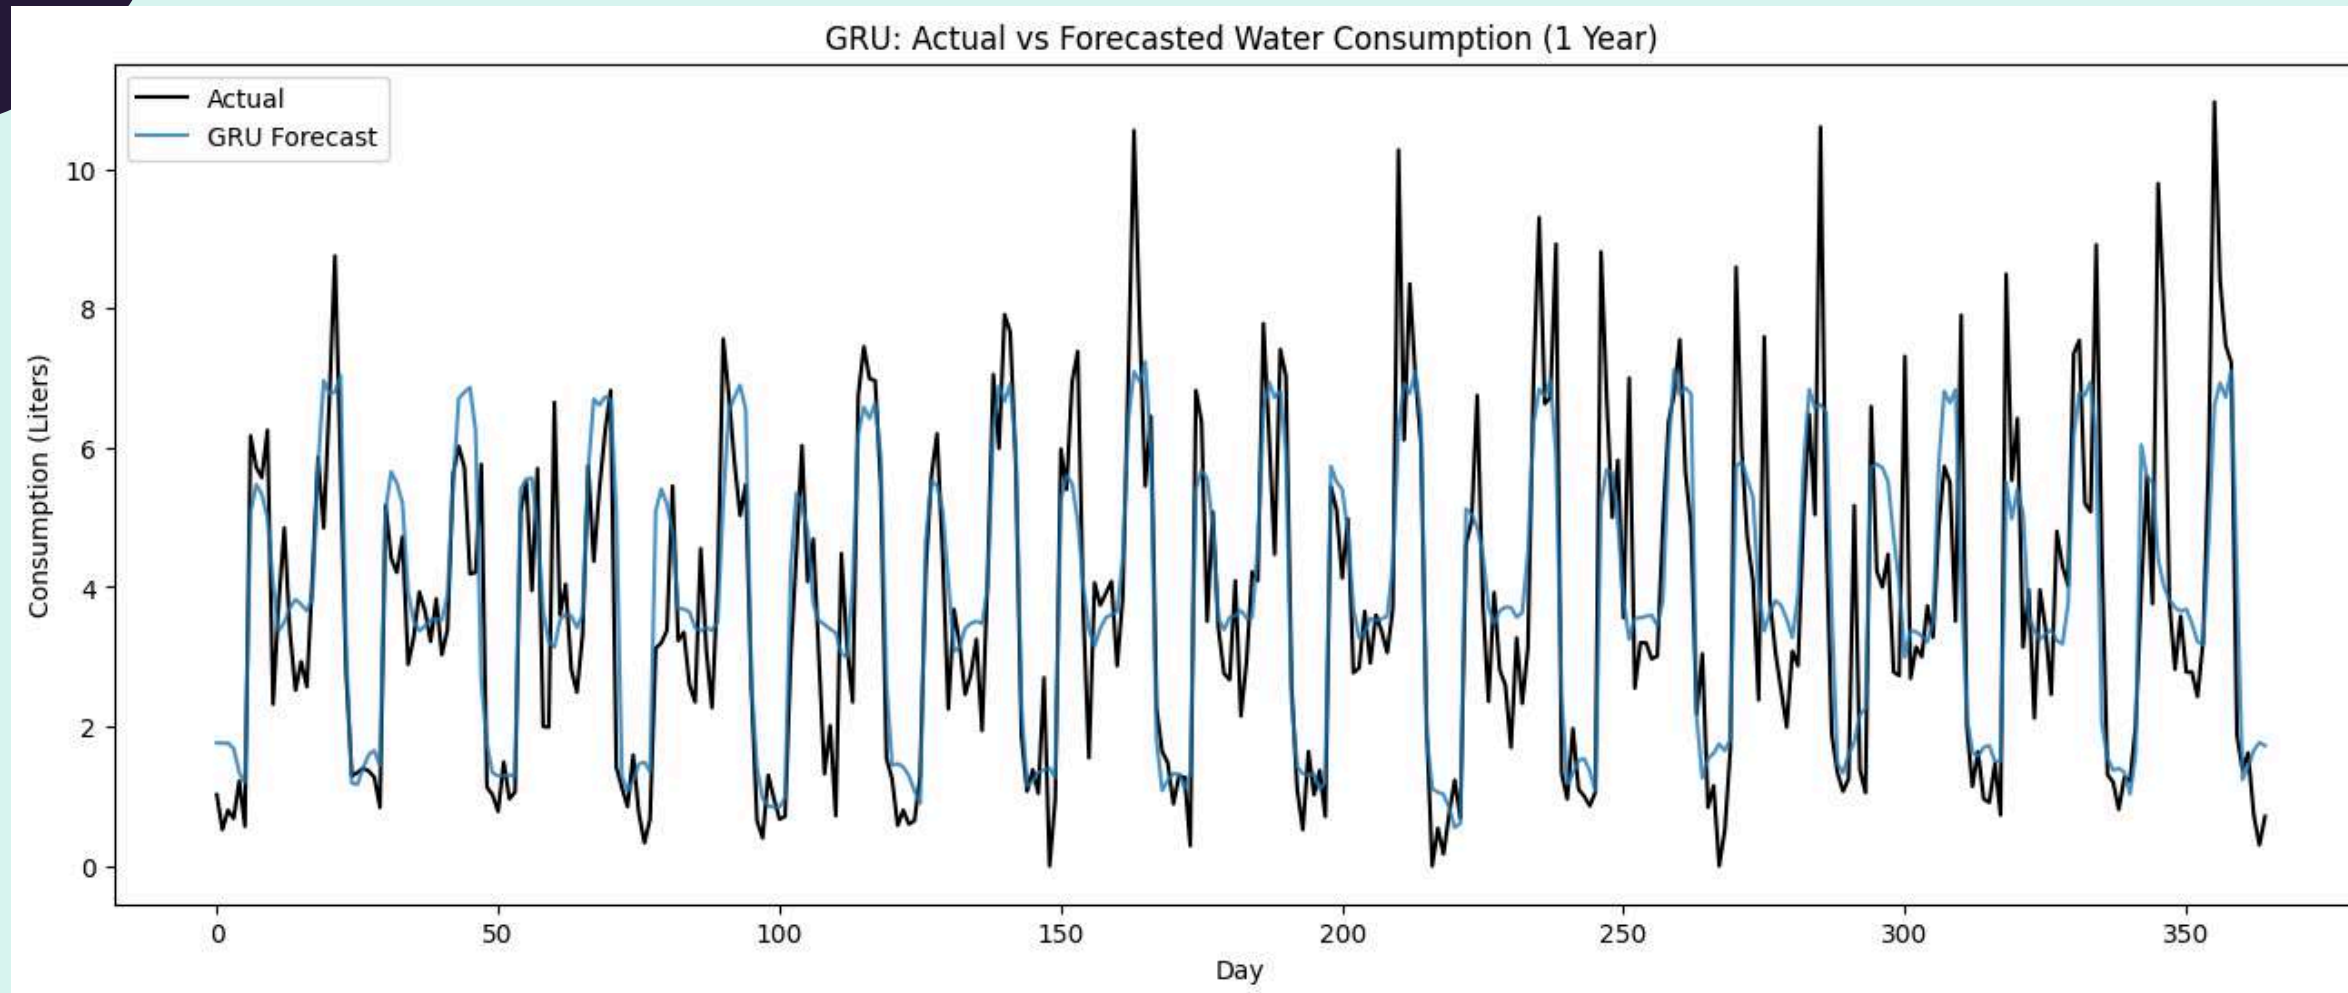


GRU

```
# Optimized GRU Model with Dropout and Early Stopping
model_gru = Sequential([
    Input(shape=(window_size, 1)),
    GRU(64, return_sequences=True),
    Dropout(0.2),
    GRU(32),
    Dropout(0.2),
    Dense(1)
])
model_gru.compile(optimizer='adam', loss='mse')
model_gru.fit(X_train, y_train, validation_split=0.1, epochs=50, batch_size=32,
              callbacks=[early_stopping], verbose=1)
```



GRU




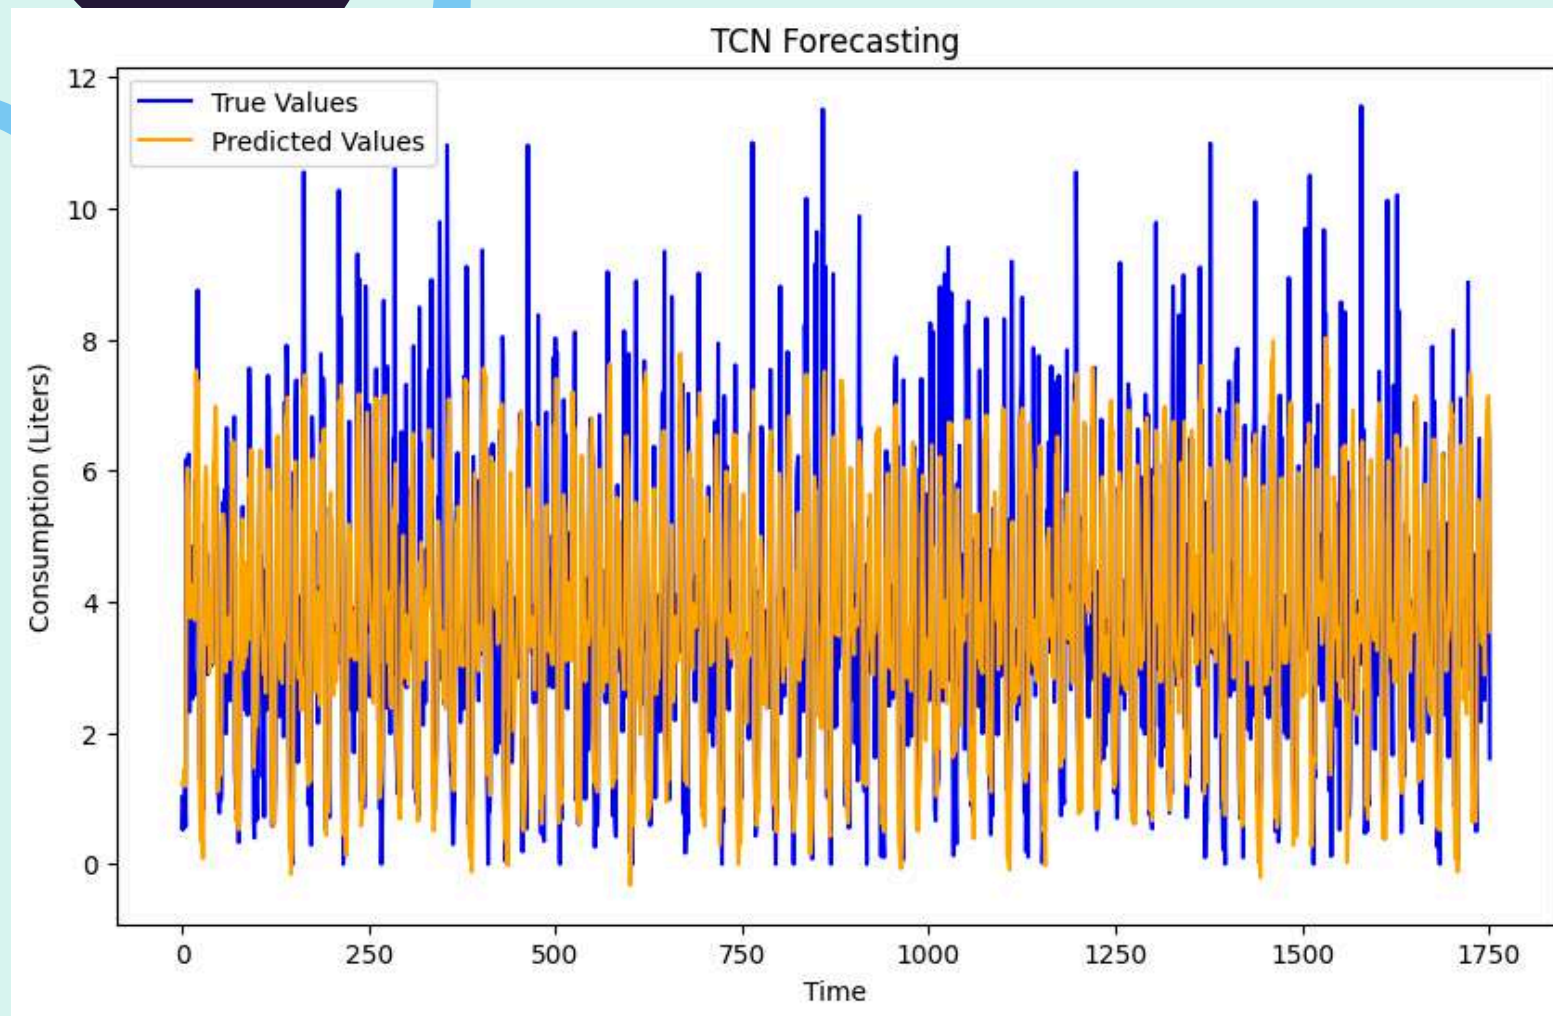
GRU (Gated Recurrent Unit)
Forecast Behavior: Similar to LSTM with stable and consistent trends, slightly more responsive to short-term variations.
Strengths: Efficient in training and performs well with fewer parameters.
Limitations: Like LSTM, may smooth out high-frequency noise or irregularities.

TCN

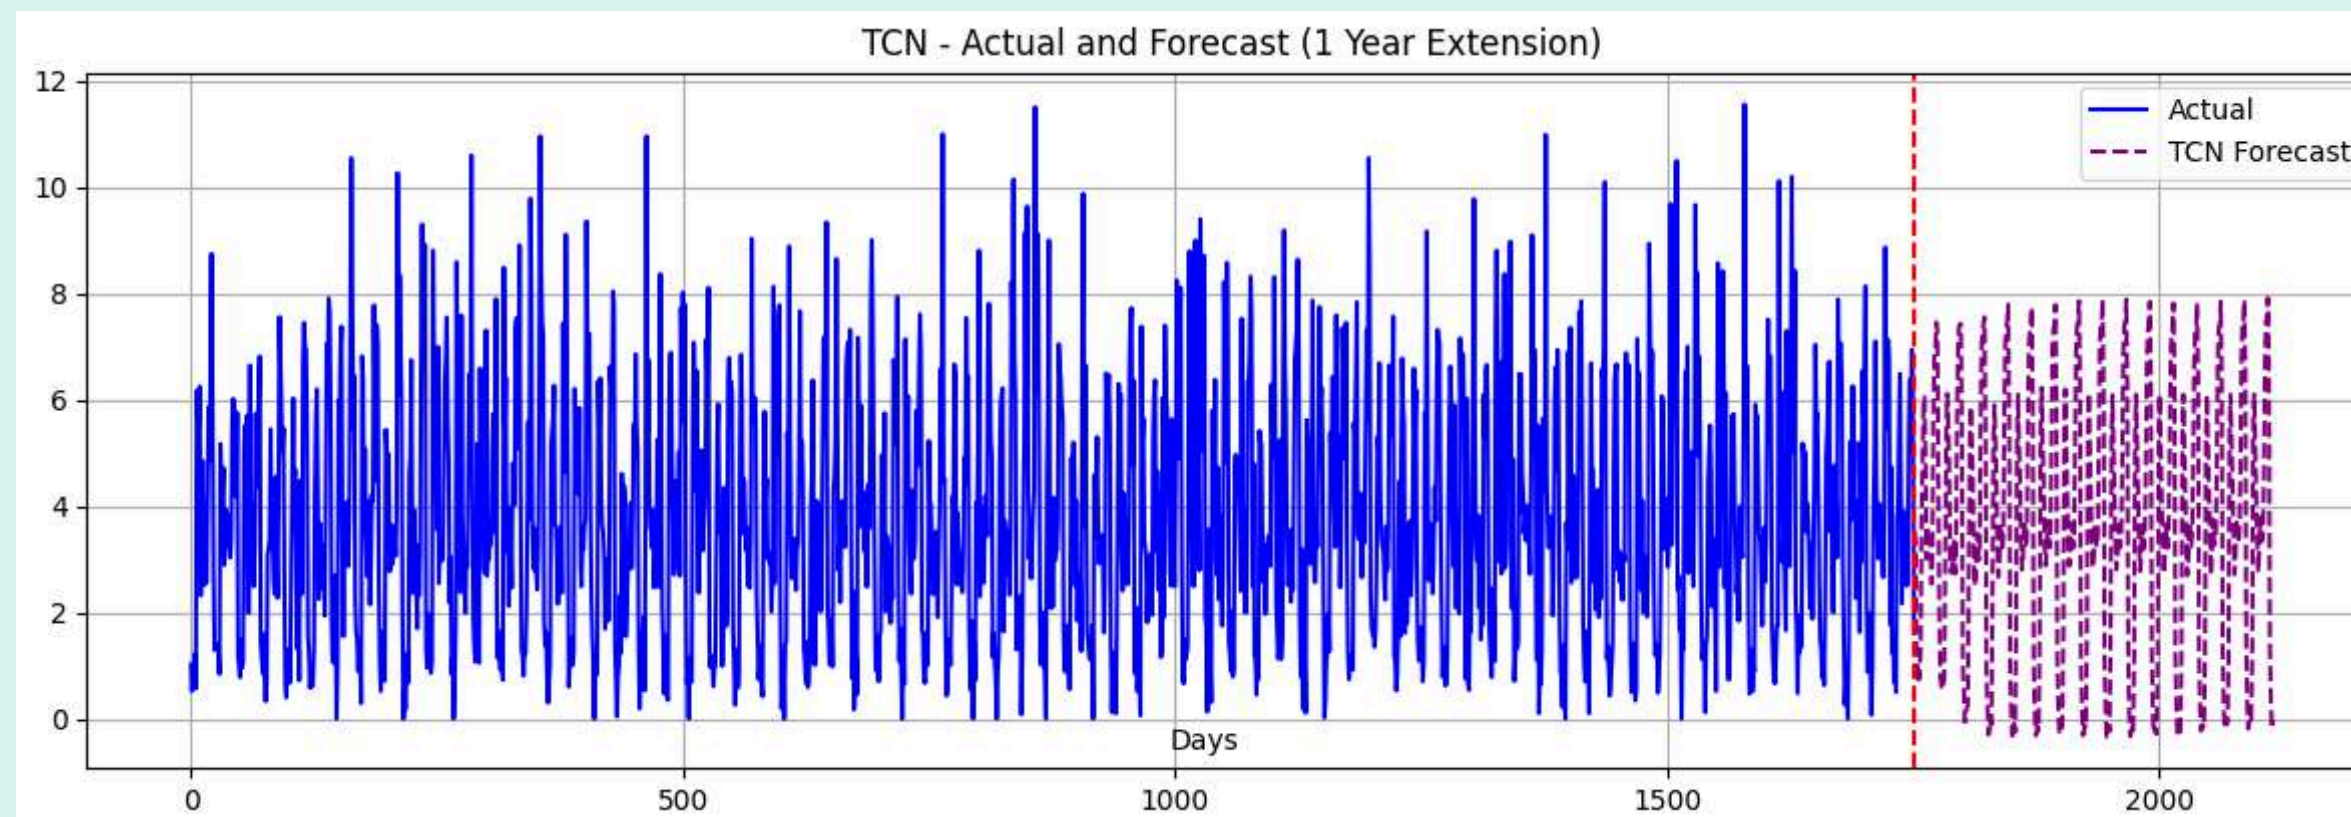
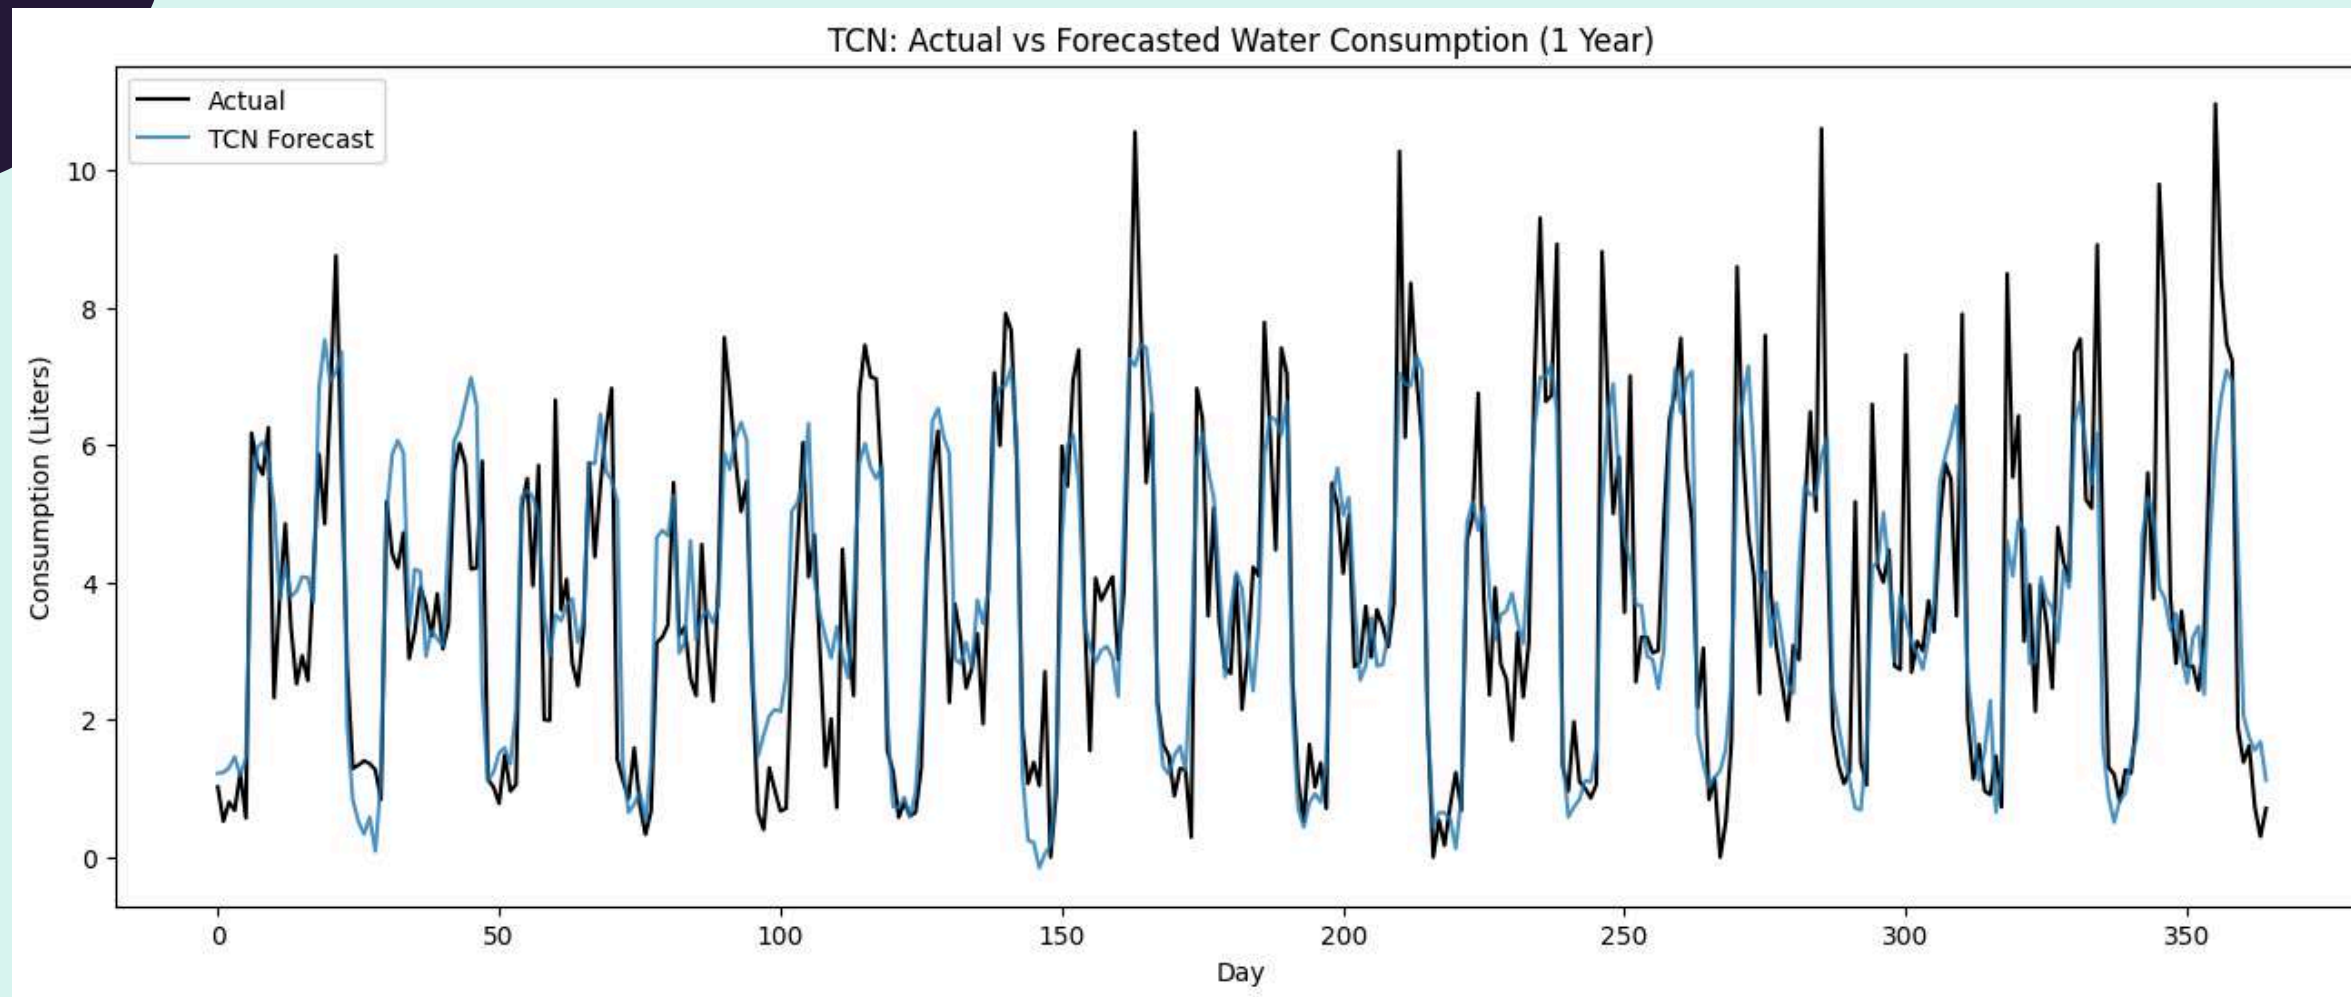
```
# Optimized TCN Model with Dropout and Early Stopping
def residual_block(x, filters, kernel_size, dilation_rate):
    shortcut = x
    x = Conv1D(filters, kernel_size, dilation_rate=dilation_rate, padding='causal')(x)
    x = BatchNormalization()(x)
    x = Activation('relu')(x)
    x = Dropout(0.2)(x)
    x = Conv1D(filters, kernel_size, dilation_rate=dilation_rate, padding='causal')(x)
    x = BatchNormalization()(x)
    if shortcut.shape[-1] != x.shape[-1]:
        shortcut = Conv1D(filters, 1, padding='same')(shortcut)
    x = Add()([shortcut, x])
    return Activation('relu')(x)

inputs = Input(shape=(window_size, 1))
x = residual_block(inputs, 32, 3, 1)
x = residual_block(x, 32, 3, 2)
x = residual_block(x, 32, 3, 4)
x = GlobalAveragePooling1D()(x)
outputs = Dense(1)(x)

model_tcn = Model(inputs, outputs)
model_tcn.compile(optimizer='adam', loss='mse')
model_tcn.fit(X_train, y_train, validation_split=0.1, epochs=50, batch_size=32,
              callbacks=[early_stopping], verbose=1)
```




TCN




TCN (Temporal Convolutional Network)
Forecast Behavior: Forecast is more detailed with visible high-frequency fluctuations.
Strengths: Better at capturing short-term patterns and local variations in consumption.
Limitations: Slightly more complex and potentially more sensitive to noise.


MODELS MATRIX



The LSTM model is better at identifying Normal consumption (887 correct predictions) compared to High consumption (630 correct predictions). It misclassifies 87 instances of Normal as High and 148 instances of High as Normal.



The GRU model shows balanced performance, with slightly fewer correct predictions for Normal (848) than LSTM but better accuracy for High (645). It has fewer misclassifications for High (133) compared to LSTM, suggesting improved sensitivity for detecting High consumption.




The TCN model excels in identifying High consumption (660 correct predictions, the highest among the three models).

However, it has the lowest accuracy for Normal (822) and the highest False High errors (152), indicating it may be more conservative in labeling consumption as High.



Model	MAE ↓	MSE ↓	RMSE ↓	Performance
LSTM	0.92	1.66	1.29	★ Best
GRU	0.94	1.69	1.3	👍 Good
TCN	0.99	1.88	1.37	👎 Worst





Thank you for listening!

EMAIL

qjladornado01@tip.edu.ph

qvabapuyan01@tip.edu.ph